

Parallel Load Balancing

Minke Zhang

February 11, 2014

1 Modules

Locks

We wish to abstract away the specifics of the lock interface so that at lock invocation, we do not need to worry about how to handle each type of lock. As such, we will allow a lock struct to deal with the complexities of each lock, and to provide a standard interface for the experiments.

```
struct lock_t {
    int type;           // the type of lock being used
    int status;         // a copy of the lock status
    void *l;           // the lock
} lock;

lock *init_lock(int type);
void lock(lock *l) {
    switch(l->type) {
        ...
    }
    l->status = 1;
}

int try_lock(lock *l) {
    return(status);
}

void unlock(lock *l) {
    switch(l->type) {
        ...
    }
    l->status = 0;
}
```

Experiment Interfaces

What follows is the expected interface for experiments– the hope being that if all else is sufficient, we will simply alter the given variables and

```
int time_counter_serial(int M);           // returns the counter value at
```

```

int time_counter_parallel(int M, int n, int L);           // M seconds
                                                         // returns the counter value at
                                                         // M seconds, with n threads
                                                         // and lock type L

float work_counter_serial(int B);                       // returns the time taken for
                                                         // the counter to reach B

float work_counter_parallel(int B, int n, int L);        // returns the time taken for
                                                         // the counter to reach B,
                                                         // with n threads and lock
                                                         // type L

void packet_serial(int M, int n, long W, int            // same as the serial version
    uniform_flag, short experiment_number);            // of the last project, but
                                                         // without specifying T

void packet_parallel(int M, int n, long W, int          // parallel version, with (non)-
    uniform_flag, short experiment_number, int D,      // locking dequeues and load
    int L, int S);                                     // balancing strategy S

```

2 Hypotheses

TaS, exponential, mutex, anderson, CLH

Idle Overhead

We would expect that the overhead generated by the TAS and backoff lock will not significantly slow down the parallel counter— the length of the uncontested path is a simple straight walk through the `lock()` function, whereas the queue-based locks will require memory allocation per lock call, and as such, should incur a higher overhead penalty. The POSIX lock implementation is, as per StackOverflow (<http://bit.ly/MHnDdX>), a userspace increment-and-test implementation, which is hardware instruction based; as such, we would expect the performance of the POSIX implementation to lie closer to TAS at $n = 1$ than to the queue-based locks.

Lock Scaling

As was mentioned in class, we would expect that the TAS and backoff locks to behave rather poorly under heavy contention, whereas the queue-based locks are much better at adjusting to the heavier congestion, as the queue-based locks avoid in large part the expensive cache misses which occur for conventional locks.

Fairness

We would expect that the TAS to have a much higher deviation than the queue-based locks, due to possibilities of starvation of certain threads in cases of heavy contention, whereas the queue-based locks should have a small deviation due to the way in which lock requests are being made.

Packet Overhead

We believe there will be more overhead generated due to the homequeue lock implementation—because the nature of locks, we are serializing a process which was before perfectly fine being modified concurrently. Thus, we would expect the worker rate to be lower for implementations of the firewall utilizing locks than using the lock-free queue. Moreover, we would expect that the relative overhead between lock implementations to mirror that of the **idle overhead**.

Packet Scaling

We do not believe there to be a significant speedup while using locks in the uniform packet case—the use of locks for the firewall is to provide a way for multiple threads to call `dequeue()` on a queue simultaneously. This is only useful for the case of load-balancing, which will only be a major consideration in the exponentially distributed packet case. We will therefore expect the *awesome* strategy to emphasize high number of threads using the queue-based lock, and with some work-stealing scheduling implemented.

3 Testing

The crux of the issue here, assuming the firewall was sufficiently tested in the last project, is the ability for each lock to serialize some set of operations. The simplest— and perhaps the most effective— method for testing this is to spawn n threads making m incrementing calls each on a shared resource, and test to see at the end if the shared resource has been altered $n \times m$ times.